

Building Kafka-based Microservices with Akka Streams and Kafka Streams

Boris Lublinsky and Dean Wampler, Lightbend

boris.lublinsky@lightbend.com

dean.wampler@lightbend.com

The background features two O'Reilly book covers. The top cover is red with a black spine and the text 'Dean Wampler' in white. The bottom cover is red with a black spine and the text 'Fast Data Architectures for Streaming Applications' in white. The O'Reilly logo is visible on the spines of both books.

O'REILLY®

Fast Data Architectures for Streaming Applications


Getting Answers Now from Data Sets that Never End

By Dean Wampler, Ph. D., VP of Fast Data Engineering

Get Your Free Copy

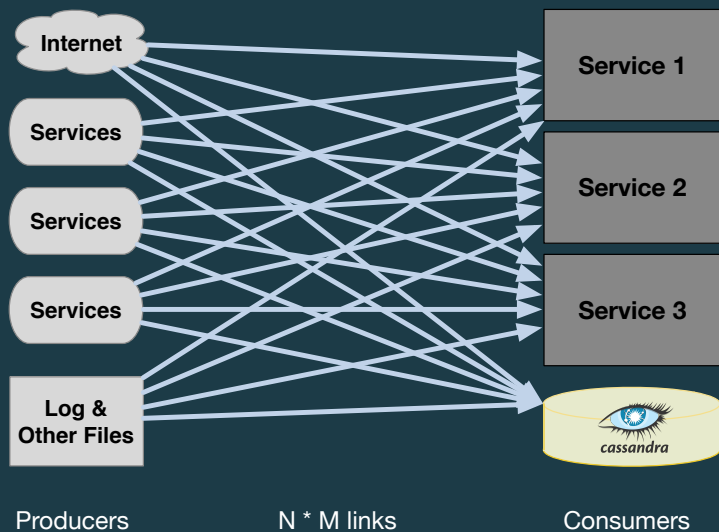
Today's focus:

- Kafka - the data backplane
- Akka Streams and Kafka Streams - streaming microservices



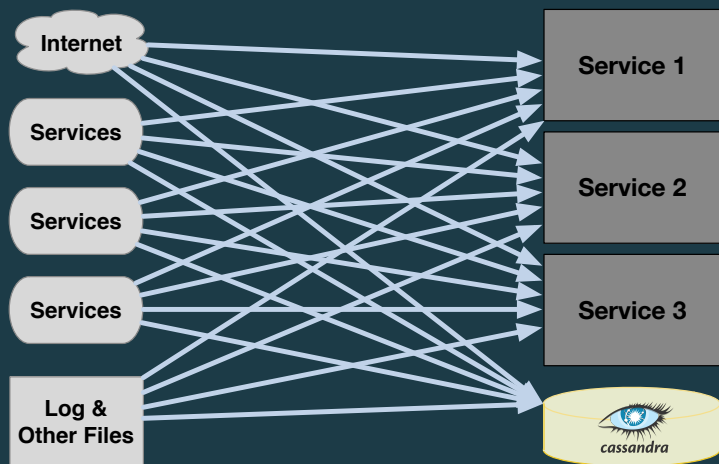
Why Kafka for Connectivity?

Before:



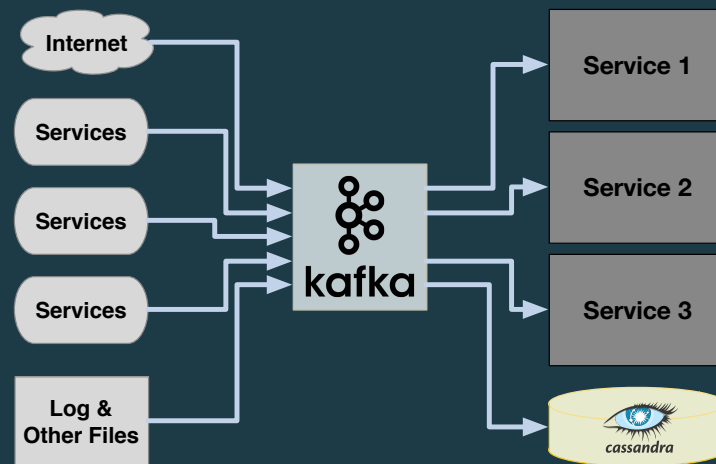
Why Kafka for Connectivity?

Before:



$N * M$ links

After:

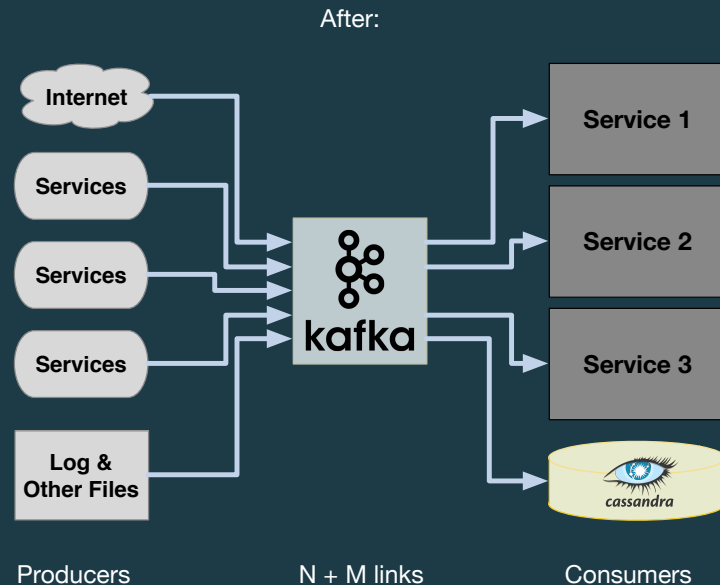


$N + M$ links

Why Kafka for Connectivity?


Kafka:

- Simplify dependencies between services
 - Improved data consistency
- Minimize data transmissions
- Reduce data loss when a service crashes
- M producers, N consumers
 - Improved extensibility
- Simplicity of one “API” for communication




Streaming Engines:

Spark, Flink - services to which you submit work. Large scale, automatic data partitioning.




Streaming Engines:

Spark, Flink - services to which you submit work. Large scale, automatic data partitioning.




Streaming Frameworks:

Akka Streams, Kafka Streams - libraries/Frameworks for “data-centric micro services”. Smaller scale, but great flexibility




A Spectrum of Microservices

Event-driven μ -services



"Record-centric" μ -services



Events

Records

A Spectrum of Microservices



Event-driven μ -services



Akka emerged from the left-hand side of the spectrum, the world of highly *Reactive* microservices.

Akka Streams pushes to the right, more data-centric.


A Spectrum of Microservices



Emerged from the right-hand side.

Kafka Streams pushes to the left, supporting many event-processing scenarios.

“Record-centric” μ -services



Machine Learning and Model Serving: A Quick introduction



O'REILLY®

Serving Machine Learning Models

**A Guide to Architecture, Stream Processing Engines,
and Frameworks**

By Boris Lublinsky, Fast Data Platform Architect

Get Your Free Copy

ML Is Simple



Data



Magic



Happiness

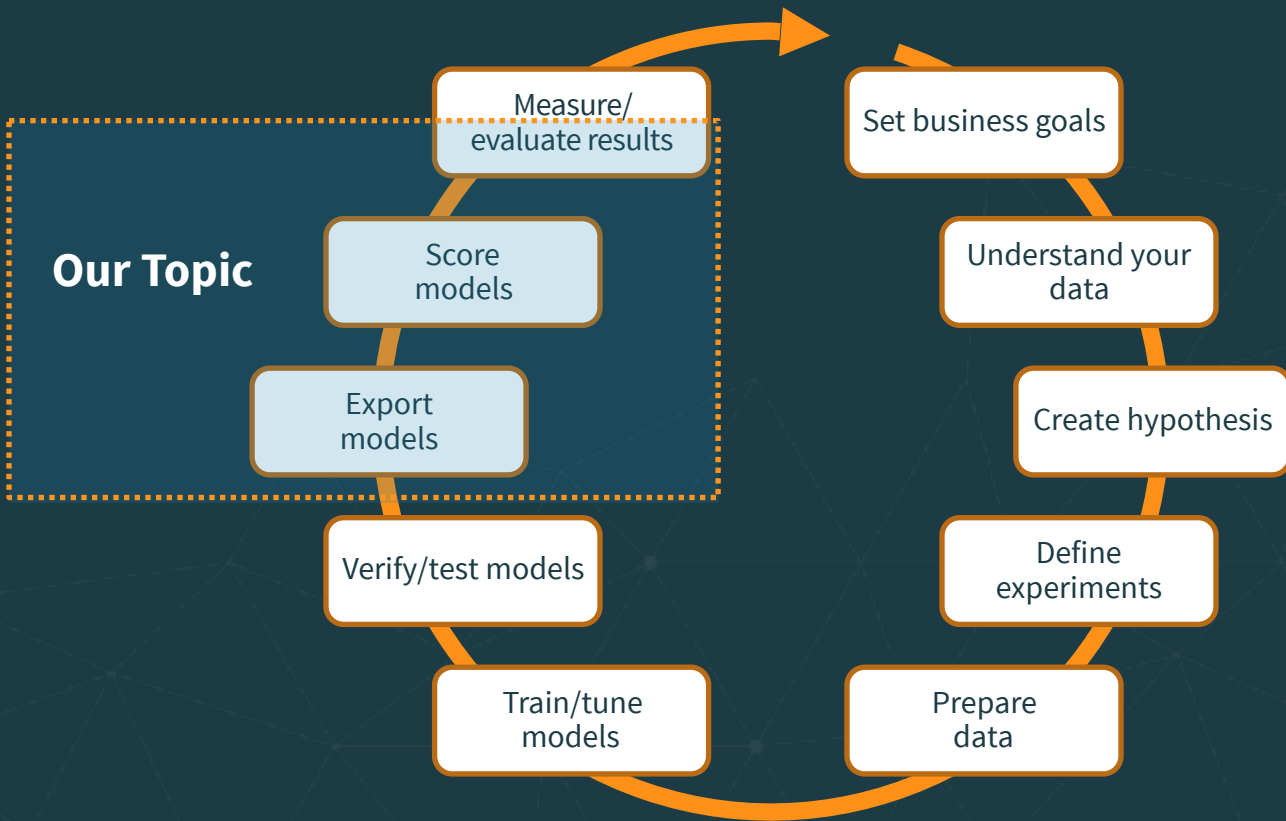
Maybe Not



Even If There Are Instructions



The Reality



What Is The Model?

A model is a function transforming inputs to outputs - $y = f(x)$

for example:

Linear regression: $y = a_c + a_1 * x_1 + \dots + a_n * x_n$


Neural network: $f(x) = K(\sum_i w_i g_i(x))$

Such a definition of the model allows for an easy implementation of model's composition. From the implementation point of view it is just function composition



Model Learning Pipeline

UC Berkeley AMPLab introduced [machine learning pipelines](#) as a graph defining the complete chain of data transformation.




Traditional Approach To Model Serving


- Model is code
- This code has to be saved and then somehow imported into model serving

Why is this problematic?

Impedance Mismatch




**Continually expanding
Data Scientist toolbox**



**Defined Software
Engineer toolbox**

Alternative - Model As Data



Standards



Portable
Format for
Analytics (PFA)



Exporting Model As Data With PMML

There are already a lot of export options



<https://github.com/jpmml/jpmml-sparkml>



<https://github.com/jpmml/jpmml-sklearn>



<https://github.com/jpmml/jpmml-r>



<https://github.com/jpmml/jpmml-tensorflow>



Evaluating PMML Model

There are also a couple PMML evaluators



<https://github.com/jpmml/jpmml-evaluator>



<https://github.com/opendatagroup/augustus>

Exporting Model As Data With Tensorflow

- Tensorflow execution is based on Tensors and Graphs
- Tensors are defined as multilinear functions which consists of various vector variables
- A computational graph is a series of Tensorflow operations arranged into graph of nodes.
- Tensorflow support exporting of such graph in the form of binary protocol buffers.
- There are two different export format - optimized graph and a new format - saved model



Evaluating Tensorflow Model

- Tensorflow is implemented in C++ with Python interface.
- In order to simplify Tensorflow usage from Java, in 2017 Google introduced Tensorflow Java APIs.
- Tensorflow Java APIs supports import of the exported model and allows to use them for scoring.



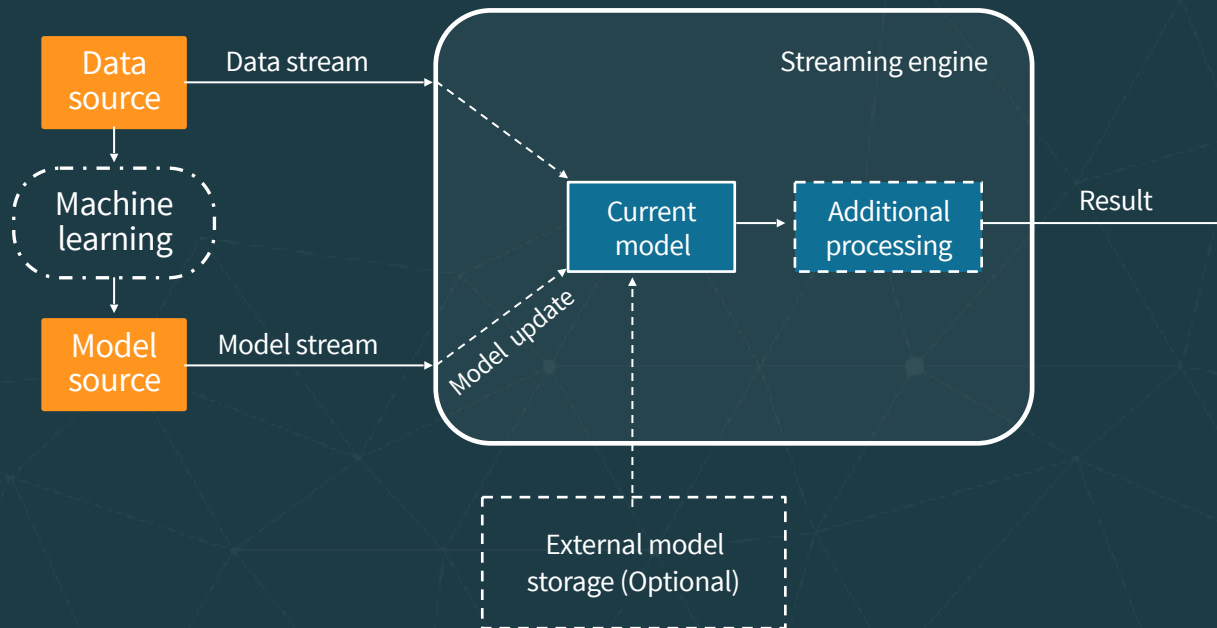
Additional Considerations – Model Lifecycle

- Models tend to change
- Update frequencies vary greatly – from hourly to quarterly/yearly
- Model version tracking
- Model release practices
- Model update process



The Solution

A streaming system allowing to update models without interruption of execution
(dynamically controlled stream).



Model Representation (Protobufs)

// On the wire

syntax = "proto3";

// Description of the trained model.

message ModelDescriptor {

string name = 1; // Model name

string description = 2; // Human readable

string dataType = 3; // Data type for which this model is applied.

enum ModelType { // Model type

TENSORFLOW = 0;

TENSORFLOWSAVED = 2;

PMML = 2;

};

ModelType modeltype = 4;

oneof MessageContent {

// Byte array containing the model

bytes data = 5;

string location = 6;

}

}

Model Representation (Scala)

```
// Internal
```

```
trait Model {  
  def score(input : AnyVal) : AnyVal  
  def cleanup() : Unit  
  def toBytes() : Array[Byte]  
  def getType : Long  
}  
  
def ModelFactoryI {  
  def create(input : ModelDescriptor) : Model  
  def restore(bytes : Array[Byte]) : Model  
}
```


Additional Considerations: Monitoring

Model monitoring should provide information about usage, behavior, performance and lifecycle of the deployed models

```
case class ModelToServeStats(  
  name: String,                // Model name  
  description: String,         // Model descriptor  
  modelType: ModelDescriptor.ModelType, // Model type  
  since : Long,                // Start time of model usage  
  var usage : Long = 0,        // Number of servings  
  var duration : Double = .0,  // Time spent on serving  
  var min : Long = Long.MaxValue, // Min serving time  
  var max : Long = Long.MinValue // Max serving time  
)
```


Queryable State

Queryable state (interactive queries) is an approach, which allows to get more from streaming than just the processing of data. This feature allows to treat the stream processing layer as a lightweight embedded database and, more concretely, to *directly query the current state* of a stream processing application, without needing to materialize that state to external databases or external storage first.



Implementation Options

Modern stream-processing engines (SPE) take advantage of the cluster architectures. They organize computations into a set of operators, which enables execution parallelism; different operators can run on different threads or different machines.

Stream-processing library (SPL), on the other hand, is a library, and often domain-specific language (DSL), of constructs simplifying building streaming applications.






- Akka Streams is a library focused on in process back-pressured reactive streaming.
- Can be used with Akka Cluster
- Provides a broad ecosystem of connectors to various technologies (data stores, message queues, file stores, streaming services, etc) - Alpakka
- Integrated with other Akka components, including Akka Actors, Akka HTTP, etc.
- In Akka Streams computations are written in a DSL for defining a graph, which aims to make translating graph drawings to and from code simpler.

Using Custom Stage

Create custom stage, which is a fully type-safe way to encapsulate required functionality. Your stage will provide functionality somewhat similar to a Flink low-level join




Using Custom Stage

Code time


Other Concerns

- Scale scoring with workers and routers, across a cluster
- Persist actor state with Akka Persistence
- Connect to *almost* anything with Alpakka
- *Lightbend Enterprise Suite*
- for production



Improve Scalability

Using the router actor to forward request to an individual actor responsible for processing request for a specific model type low-level join




Akka Streams with actors

Code time


Using Akka Cluster

Two level of scalability:


- Kafka partitioned topic allow to scale listeners according to the amount of partitions.
- Akka cluster sharing allows to split model serving actors across clusters.



Go Direct or Through Kafka?



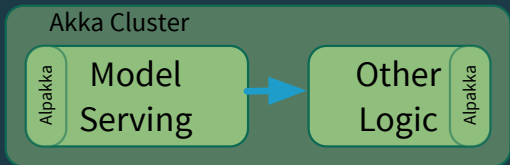
vs.



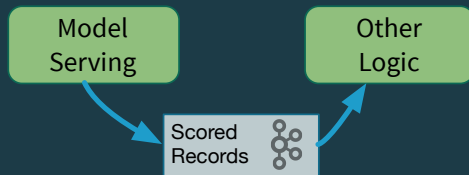
- Extremely low latency
- Minimal I/O and memory overhead
- No marshaling overhead
- Reactive Streams* back pressure
- Higher coupling - M producers, N consumers, but directly connected (sort of)
- Use Akka Persistence for durable state

- Higher latency (including queue depth)
- Higher I/O and processing (marshaling) overhead
- Better potential reusability
- Better decoupling - M producers, N consumers, completely disconnected
- Automatic durability (topics on disk)

Go Direct or Through Kafka?



VS.



- Use for smaller, faster messaging between “components”.
- Watch for consumer “backup”
- Use Akka Persistence for important state!

- Better reusability and decoupling
- Use for larger volumes, more course-grained service interactions
- Plan partitioning and replication carefully



Kafka Streams

- Kafka Streams, is a client library for processing and analyzing data stored in Kafka.
- Provides important stream-processing concepts, such as properly distinguishing between event time and processing time, windowing support, etc.
- Provides simple yet efficient management of application state based on stream/table duality.
- Provides two types of APIs:
 - Process Topology (compare to Storm)
 - DSL based on collection transformations (compare to Spark, Flink)
- Provides Java and Scala (new) APIs



Kafka Streams


- Provides low overhead
- Provides additional integrations leveraging Kafka Connect
- Provides support for scalability and load balancing (based on Kafka partitioning)
- Provides SQL interface (leveraging a separate KSQL application)
- Queryable State



Kafka Streams

- Ideally suited for:
 - ETL - leveraging KStreams
 - Aggregations - Leveraging KTable
 - Flexible integration model - in memory or over Kafka
 - Joins, including Stream and Table joins
 - Effectively once semantics

Model Serving With Kafka Streams



State Store Options

- Naive, in memory store.
- Standard key/value store provided by Kafka Streams
- Custom store

Model Serving With Kafka Streams

Code time

Thank You

Any questions?

